



**Module 4 :**

# Big Data Analysis using Apache Spark

ALY 6110: Big Data and Data Management  
Instructor: Ajit Appari

# Module 3: Topics

- Introduction to Spark
- Key Modules
- How does Spark work?
- Installation of Spark
- Load data into Spark
- Building Spark Application using Maven
  
- Reference Books for Machine Learning using Spark
  - Adi Polak (2023). [Scaling Machine Learning with Spark : distributed ML with MLlib, Tensorflow, and PyTorch](#)
  - Luu Hien (2021). [Beginning Apache Spark 3 : with DataFrame, Spark SQL, structured streaming, and Spark machine learning Library](#)
  - Jules S. Damji, Denny Lee, Brooke Wenig, Tathagata Das, Denny Lee (2020). [Learning Spark: lightning-fast data analytics](#)
  - Butch Quinto (2020) [Next-Generation Machine Learning with Spark : Covers XGBoost, LightGBM, Spark NLP, Distributed Deep Learning with Keras, and More.](#)
  - Javier Luraschi, Kevin Kuo, Edgar Ruiz (2019). [Mastering Spark with R:](#)  <https://therinspark.com/>
  - Micheal Bowles (2019). [Machine Learning with Spark and Python.](#)

# Introduction to Spark


# Introduction to Spark

- Spark is an open-source software solution that performs rapid calculations on in-memory distributed datasets referred as RDDs (Resilient Distributed Datasets);
  - RDD is a distributed collections of objects that can be cached in memory across cluster and can be manipulated in parallel;
  - RDD could be automatically recomputed on failure;
- Spark is designed as extension of MapReduce model to efficiently support a wide range of computational workloads including batch applications, iterative algorithms, interactive queries and stream processing;
- Spark is designed to be highly accessible:
  - offers simple APIs in Python, Java, Scala, SQL, and
  - Has rich built-in libraries;

# Key Modules

# Key Modules

- Spark has six core modules
- **Spark Core:** contains the basic functionality of Spark, including components:
  - Task scheduling;
  - Memory management;
  - Fault recovery;
  - Interacting with storage systems, etc.



- **Spark SQL:** is Spark's package for working with structured data
  - allows querying data via SQL as well as the Apache Hive variant of SQL—Hive Query Language (HQL)—and it supports many sources of data, including Hive tables, Parquet, and JSON
  - allows intermixing of SQL queries with the programmatic data manipulations in Python, Java, and Scala, all within a single application, thus combining SQL with complex analytics.

# Key Modules

- **Spark Streaming:** enables processing of live streams of data such as log files of web servers, or queues of messages.
  - **Spark Streaming API** for manipulating data streams closely matches the Spark Core's RDD API, making it easy to move between apps to manipulate data in memory, on disk, or arriving in real time;
- **MLib:** is a library containing common machine learning functionalities
  - Provides multiple types of machine learning algorithms, including classification, regression, clustering, and collaborative filtering, as well as supporting functionality such as model evaluation and data import;
  - All of ML methods are designed to scale out across a cluster.
- **GraphX:** is a library for manipulating graphs (e.g., a social network's friend graph) and performing graph-parallel computations;
  - Extends the Spark RDD API, allowing to create a directed graph with arbitrary properties attached to each vertex and edge;
  - Provides various operators for manipulating graphs (e.g., subgraph and mapVertices) and a library of common graph algorithms (e.g., PageRank and triangle counting).

# Key Modules

- **Cluster Manager:** allow Spark to efficiently scale up from one to many thousands of compute nodes;
  - Spark can run over a variety of cluster managers, including
    - Hadoop YARN,
    - Apache Mesos,
    - Standalone Scheduler and a simple cluster manager included in Spark itself);

# How does Spark work?

# How does Spark work?

- RDD: Your data is loaded in parallel into structured collections
- Actions: Manipulate the state of the working model by forming new RDDs and performing calculations upon them
- Persistence: Long-term storage of an RDD's state
- Spark Application is a definition in code of
  - RDD creation
  - Actions
  - Persistence
- Spark Application results in the creation of a DAG (Directed Acyclic Graph)
  - Each DAG is compiled into stages
  - Each Stage is executed as a series of Tasks
  - Each Task operates in parallel on assigned partitions
  - It all starts with the [SparkContext 'sc'](#)



# How does Spark work?

**Spark is much smarter, it keeps all results in memory**



**Distributed memory is 10-100X faster than network and disk Data Sharing in Spark**

Spark's speed comes from its ability to run computations in memory;



# Installation of Spark

# Installation of Spark

- Visit <http://spark.apache.org/downloads> and fetch the latest release of Spark;
- There are a few options on what can be downloaded:
  - The source code;
  - Binaries;



- Use Maven for Source Code installation



How to install & configure Apache Maven

The screenshot shows the "Download Apache Spark™" section of the website. It includes a numbered list of steps: 1. Choose a Spark release: 2.3.1 (Jun 08 2018) ▾; 2. Choose a package type: Pre-built for Apache Hadoop 2.7 and later ▾; 3. Download Spark: spark-2.3.1-bin-hadoop2.7.tgz; 4. Verify this release using the 2.3.1 signatures and checksums and project release KEYS. A note at the bottom states: Note: Starting version 2.0, Spark is built with Scala 2.11 by default. Scala 2.10 users should download the Spark source package and build with Scala 2.10 support.

## Link with Spark

Spark artifacts are hosted in [Maven Central](#). You can add a Maven dependency with the following coordinates:

```
groupId: org.apache.spark  
artifactId: spark-core_2.11  
version: 2.3.1
```

## Installing with PyPi

PySpark is now available in pypi. To install just run `pip install pyspark`.

## Release Notes for Stable Releases

- Spark 2.3.1 (Jun 08 2018)
- Spark 2.3.0 (Feb 28 2018)
- [Spark 2.2.2 / Jul 02 2018](#)

### Latest News

Spark+AI Summit (October 2-4th, 2018, London) agenda posted (Jul 24, 2018)

Spark 2.2.2 released (Jul 02, 2018)

Spark 2.1.3 released (Jun 29, 2018)

Spark 2.3.1 released (Jun 08, 2018)

[Archive](#)



[Download Spark](#)

### Built-in Libraries:

SQL and DataFrames

Spark Streaming

MLlib (machine learning)

GraphX (graph)

Third-Party Projects

# Installation of Spark

- Working with Cloudera's VM-s we will do the following:

- Make sure we have JDK 7+ installed;
  - Make sure we have hadoop-client package installed;

- For example, you could do:

`$ sudo yum install hadoop-client`

- `yum` will either install the package or tell you it is already there;

- Make sure all HDFS and YARN services are installed and running;

- Stop them and start them all;

- Use `yum` to run installation of Spark, by typing all on one line:

`$ sudo yum install spark-core spark-master spark-worker spark-history-server  
spark-python`

- Answer `Yes` to all the questions during installation process executed by `yum`

# Installation of Spark

- Set your SPARK\_HOME environment variable to `/usr/lib/spark`
- Add `$SPARK_HOME/bin` to your PATH environmental variable.
  - .... And you are all set now!
- `yum` made sure that you have Spark version which matches your Hadoop (Yarn) version;
- Without any further changes we will run Spark locally on this single machine;

# Spark Interactive Shells for Python, Scala, R

- Spark comes with interactive shells that enable ad hoc data analysis.
- Spark's shells allow interaction with data distributed on disks or in memory across many machines.
- Spark can load data into memory on the worker nodes, and distributed computations, even ones that process large volumes of data across many machines, can run in a few seconds;
- This makes iterative, ad hoc, and exploratory analysis commonly done in shells a good fit for Spark;


# Spark Interactive Shells for Python, Scala, R

- Spark provides both (and only) Python and Scala shells that have been augmented to support access to a cluster of machines.
- Python shell opens with [PySpark](#) command;
- Scala shell is very similar and opens with `spark-shell` command;
- You can program Spark from R using package [SparkR](#);

# Spark Interactive Shells: PySpark

## PySpark

```
[cloudera@localhost conf]$ pyspark
Python 2.8.1 (r231:90122, Aug 23 2025, 23:20:16)
[GCC 4.4.7 20201212 (Red Hat 7.4.1-11)] on linux2
Type "help", "copyright", "credits" or "license" for more information. SLF4J: Class path contains multiple
SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/zookeeper/lib/slf4j-log4j12-
1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/lib/flume-ng/lib/slf4j-log4j12-
1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple\_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory] Welcome to
```



Using Python version 2.6.6 (r266:84292, Jul 23 2015 15:22:56)  
SparkContext available as sc, HiveContext available as sqlContext.

>>>

# Spark Interactive Shells: PySpark

- The output is long and annoying. It would have been even longer had we not created log4j.properties file in the directory \$SPARK\_HOME/conf;
- You create that file by copying provided file log4j.properties.template and by changing line:
  - log4j.rootCategory=INFO, console to: log4j.rootCategory=ERROR, console
  - That lowered the logging level so that we see only the ERROR messages, and above;
- Another option is WARN, which is more verbose than ERROR but less than INFO;
- Before we proceed, let's see what kind of files and how many lines we have in HDFS
- `hadoop fs -ls ulysis`

-rw-r--r-- 1 cloudera 5258688 2015-04-01 14:32 input/4300.txt

- `hadoop fs -cat ulysis/4300.txt | wc`

33056 267980 1573079

# Load data into Spark

# Load data into Spark

- **Load Data (RDD) from HDFS {Python commands}**
- We create an RDD when we load some data (i.e. a file) into a shell variable:
- `>>> lines = sc.textFile("ulyssis/4300.txt")`
- `>>> lines.count()`  
33056
- We populated that RDD called “`lines`” with data in HDFS file “`ulyssis/4300.txt`”;
  - “`sc`” stands for an implicit `SparkContext` that allows us to communicate with the execution environment;
  - `RDD`-s are objects and has many methods,
  - The method `count()` gives the number of lines in file `4300.txt`;

# Load data into Spark

- **Load Data (RDD) from Local File (i.e. on your VM desktop) {Python commands}**
- You can find the 4300.txt file in **/home/cloudera** directory and could do the following:
  - We create an RDD when we load some data (i.e. a file) into a shell variable:
  - `>>> blines = sc.textFile("file:///home/cloudera/4300.txt")`
  - `>>> blines.count()`  
33056
- You can print the first row using method
- `>>> blines.first()`

# Load data into Spark: Source Formats

- Spark makes it very simple to load and save data in a large number of file formats;
- Spark transparently handles compressed formats based on the file extension;
- We can use both Hadoop's new and old file APIs for keyed (or paired) data;
- We can use those only with key/value data;

Format Name	Structured	Comments
Text File	No	Plain old text files. Records assumed to be one per line.
JSON	Semi	Common text-based format, semi-structured; most libraries require one record per line.
CSV	YES	Very common text-based format, often used with spreadsheet applications.
SequenceFile	YES	A common Hadoop file format used for key/value data
Protocol buffer	YES	A fast, space-efficient multi-language format
Object file	YES	Useful for saving data from a Spark job to be consumed by shared code. Breaks if you change your classes, as it relies on Java Serialization.